

# Group project (Advance Data mining and Predictive Analytics)

## - Group-8

2023-05-04

#Loading necessary poackages for current project:

```
library(caret)
## Loading required package: ggplot2
## Loading required package: lattice
library(glmnet)
## Loading required package: Matrix
## Loaded glmnet 4.1-7
library(pls)
##
## Attaching package: 'pls'
## The following object is masked from 'package:caret':
##
##      R2
## The following object is masked from 'package:stats':
##
##      loadings
library(dplyr)
##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##      filter, lag
## The following objects are masked from 'package:base':
##
##      intersect, setdiff, setequal, union
library(esquisse)
library(ggplot2)
library(randomForest)
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:dplyr':
##
##      combine
##
## The following object is masked from 'package:ggplot2':
##
##      margin
```

#Loading train dataset:

```
bank_model<-read.csv("train_v3.csv")
```

#Creating a default column based on, if loss is 0 then 0 and if loss is more than 0 then default is 1

```
bank_model$default <- factor(ifelse(bank_model$loss > 0, 1, 0))
```

```
bank_model$loss <- (bank_model$loss / 100)
```

#Checking the missing values in the data set:

```
row_missing <- rowMeans(is.na(bank_model))
```

```
min_missing_values<-min(row_missing)
min_missing_values
```

```
## [1] 0
```


```
max_missing_values<-max(row_missing)
max_missing_values
```


```
## [1] 0.4790576
```

**we have maximum missing values in the dataset with percentage of 47%**

## Visualizations of the dataset:

```
ggplot(bank_model, aes(x=factor(default))) +
  geom_bar(stat="count", width=0.4, fill="blue") +
  labs(title="Non-Default v/s Default") +
  labs(x="", y="No. of Customers") +
  theme(plot.title = element_text(hjust = 0.4)) +
  geom_text(stat='count', aes(label=..count..), vjust=2)
```

```
## Warning: The dot-dot notation (`..count..`) was deprecated in ggplot2
3.4.0.
##  Please use `after_stat(count)` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```



#Removing the zero-variances variables and Preprocessing the dataset by removing highly correlated and imputing missing values using “corr” and “medianimpute”:

```
zero_var_indices <- nearZeroVar(bank_model[, -c(763,764)])
data_cleaned <- bank_model[, -zero_var_indices]
bank_preprocess <- preProcess(data_cleaned[, -c(739,740)], method = c("corr",
"medianImpute"))
new_bank_model <- predict(bank_preprocess, data_cleaned)
```

#Removing zero-variance, highly corr variables and imputed missing values : we have new data set “new\_bank\_model” with 248 attributes.

#1.CLASSIFICATION MODEL:

#We first need to run a classification model to classify how many customers are actually defaulting

#We used Lasso and Principle Component Analysis(PCA) for variable selection

#a).Lasso Model: we now run the new\_bank\_model with 248 attributes for variable selection:


```
set.seed(123)

y <- as.vector(as.factor(new_bank_model$default))

x <- data.matrix(new_bank_model[, -c(247, 248)])

lasso_model <- cv.glmnet(x, y, alpha = 1, preProcess =
c("center", "scale"), family = "binomial", nfolds = 10, type.measure = "auc")

plot(lasso_model)
```



```
lasso_model$lambda.min
```

```
## [1] 0.0002640483
```

#Minimum Lambda value returned a total of 180 attributes out of 248 attributes:

#We convert coefficients returned in lasso model into dataframe:

```
# Return the coefficients for the lasso regression at the minimum lambda value:
```

```
coef <- coef(lasso_model, s= "lambda.min")
```

```
#Convert the coefficient values into a dataframe:
```

```
new_bank_coef<- data.frame(name = coef@Dimnames[[1]][coef@i + 1], coefficient = coef@x)
```

```
#Removing negatives values using "abs" function:
```

```
new_bank_coef$coefficient <- abs(new_bank_coef$coefficient)
```

```
#Re-arranging the data frame in decreasing order:
```

```
new_bank_coef[order(new_bank_coef$coefficient, decreasing = TRUE), ]
```

```
##           name coefficient
## 1  (Intercept) 7.211311e+00
## 28          f129 2.418453e+00
## 59          f268 1.679735e+00
## 116         f471 1.051069e+00
## 178         f768 1.034926e+00
## 13           f57 8.977787e-01
## 136         f604 8.537344e-01
## 25           f99 6.190194e-01
## 71          f298 4.921528e-01
## 74          f306 4.318747e-01
## 55          f240 3.932927e-01
## 128         f546 3.818846e-01
## 177         f765 3.700267e-01
## 127         f536 3.259373e-01
## 179         f774 3.232158e-01
## 37          f153 3.131973e-01
## 140         f615 2.945556e-01
## 176         f756 2.802011e-01
## 18           f70 2.738019e-01
## 129         f556 2.566507e-01
## 56          f243 2.504861e-01
## 145         f630 2.465019e-01
## 15          f65 2.434348e-01
## 14          f61 2.341287e-01
## 45          f198 2.208872e-01
## 70          f297 2.160104e-01
## 69          f292 2.132267e-01
## 27          f104 1.840168e-01
## 77          f313 1.644578e-01
## 118         f479 1.625639e-01
## 141         f616 1.605380e-01
## 19          f71 1.572015e-01
## 22          f80 1.526873e-01
```

## 68	f290	1.512196e-01
## 133	f590	1.404800e-01
## 23	f81	1.309131e-01
## 134	f598	1.243174e-01
## 83	f331	1.217854e-01
## 80	f321	1.189038e-01
## 143	f620	1.060489e-01
## 81	f323	1.020564e-01
## 17	f67	9.998148e-02
## 40	f170	9.622995e-02
## 96	f384	9.571386e-02
## 147	f637	8.571398e-02
## 24	f83	8.259875e-02
## 52	f218	7.780791e-02
## 86	f341	7.624593e-02
## 6	f13	7.019709e-02
## 85	f340	6.972299e-02
## 75	f308	6.938225e-02
## 112	f458	6.932195e-02
## 63	f279	6.233926e-02
## 26	f103	6.178347e-02
## 57	f252	6.047895e-02
## 54	f222	5.931565e-02
## 39	f163	5.769257e-02
## 79	f316	5.764058e-02
## 50	f213	5.644315e-02
## 126	f533	5.478935e-02
## 3	f3	5.138892e-02
## 16	f66	4.860839e-02
## 110	f448	4.642512e-02
## 67	f289	4.608358e-02
## 58	f262	4.379036e-02
## 175	f755	4.341701e-02
## 73	f304	3.884648e-02
## 53	f220	3.778266e-02
## 114	f468	3.749926e-02
## 111	f450	3.623018e-02
## 82	f330	3.499625e-02
## 109	f444	3.409257e-02
## 135	f601	3.397650e-02
## 97	f385	3.396528e-02
## 121	f518	3.283955e-02
## 148	f638	3.229011e-02
## 138	f612	3.127051e-02
## 38	f162	2.911762e-02
## 93	f374	2.614121e-02
## 132	f589	2.433474e-02
## 181	f776	2.369116e-02
## 29	f130	2.339993e-02
## 131	f588	2.293142e-02

## 180	f775	2.213799e-02
## 108	f442	2.174456e-02
## 137	f609	2.153892e-02
## 20	f73	1.965623e-02
## 89	f350	1.874173e-02
## 163	f677	1.812658e-02
## 44	f189	1.682807e-02
## 32	f143	1.606743e-02
## 60	f269	1.570063e-02
## 66	f288	1.565904e-02
## 95	f383	1.564887e-02
## 159	f669	1.551487e-02
## 151	f647	1.542845e-02
## 8	f19	1.542647e-02
## 36	f150	1.526814e-02
## 104	f422	1.457757e-02
## 11	f44	1.445819e-02
## 168	f725	1.436713e-02
## 122	f522	1.407689e-02
## 162	f674	1.309686e-02
## 92	f367	1.303513e-02
## 160	f672	1.297878e-02
## 142	f619	1.293453e-02
## 88	f349	1.228726e-02
## 169	f733	1.196646e-02
## 21	f76	1.192408e-02
## 158	f664	1.171386e-02
## 31	f140	1.165022e-02
## 12	f54	1.096574e-02
## 94	f378	1.082223e-02
## 150	f646	1.040117e-02
## 124	f525	1.014318e-02
## 72	f299	9.879784e-03
## 78	f315	8.661046e-03
## 120	f514	8.568343e-03
## 46	f199	7.107456e-03
## 41	f173	6.782461e-03
## 10	f32	5.922245e-03
## 170	f734	5.686606e-03
## 90	f358	5.336970e-03
## 101	f411	5.190485e-03
## 164	f679	5.088550e-03
## 161	f673	4.987415e-03
## 172	f739	4.861803e-03
## 100	f403	4.795423e-03
## 107	f436	4.068424e-03
## 5	f5	3.986515e-03
## 35	f149	3.900669e-03
## 105	f425	3.866593e-03
## 42	f182	3.805911e-03

## 43	f188 3.176033e-03
## 76	f312 2.384219e-03
## 33	f144 2.138026e-03
## 153	f650 2.102708e-03
## 49	f212 2.035054e-03
## 154	f651 1.666227e-03
## 149	f645 1.516828e-03
## 156	f654 1.369220e-03
## 9	f29 6.459194e-04
## 155	f652 4.109182e-04
## 165	f680 4.050047e-04
## 119	f499 3.748039e-04
## 91	f361 3.399257e-04
## 152	f649 3.317337e-04
## 139	f614 2.300847e-04
## 64	f285 5.024071e-05
## 4	f4 2.852757e-05
## 106	f428 2.392520e-05
## 173	f742 2.115615e-05
## 130	f587 1.530323e-05
## 99	f394 6.394796e-06
## 84	f333 2.550679e-06
## 62	f278 2.135087e-06
## 30	f139 1.732321e-06
## 48	f203 1.200118e-06
## 2	id 1.050045e-06
## 34	f146 2.250796e-07
## 103	f421 2.246499e-07
## 146	f636 1.233782e-07
## 47	f202 7.852488e-08
## 174	f743 3.973979e-08
## 166	f699 1.719785e-08
## 115	f470 1.164576e-08
## 123	f523 1.027928e-08
## 65	f287 3.415609e-09
## 7	f16 1.468087e-09
## 171	f735 1.305841e-09
## 167	f715 6.683286e-10
## 125	f526 5.990584e-10
## 51	f217 5.529375e-10
## 144	f628 1.143328e-10
## 87	f347 1.380599e-11
## 157	f659 5.688570e-14
## 113	f461 1.279305e-17
## 102	f420 8.661696e-26
## 61	f277 2.839730e-27
## 117	f472 9.415087e-37
## 98	f391 4.903228e-43




```

#Removing intercept columns returned from Lasso model:
new_bank_coef<- new_bank_coef[-1, ]

#Converting the data frame to a vector:
new_bank_coef<- as.vector(new_bank_coef$name)

#Adding "default" column to the data frame:
new_bank_coef<- c(new_bank_coef,"default")

#Selecting attributes from original data set "new_bank_model" using
coefficients returned from Lasso model i.e., "new_bank_coef"
bank_lasso<-select(new_bank_model, new_bank_coef)

## Warning: Using an external vector in selections was deprecated in
tidyselect 1.1.0.
##  Please use `all_of()` or `any_of()` instead.
##   # Was:
##   data %>% select(new_bank_coef)
##
##   # Now:
##   data %>% select(all_of(new_bank_coef))
##
## See <https://tidyselect.r-lib.org/reference/faq-external-vector.html>.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

```

#### #b). Principle Component Analysis (PCA):

#We have 180 variables returned from Lasso model that are stored in “bank\_lasso”. Now, we further process the variables using PCA.

```

pca_model <- preProcess(bank_lasso[, -c(181)], method = c("center", "scale",
"pca"), thresh = 0.80)

pca_model_1<- predict(pca_model, bank_lasso)

pca_model

## Created from 80000 samples and 180 variables
##
## Pre-processing:
##   - centered (180)
##   - ignored (0)
##   - principal component signal extraction (180)
##   - scaled (180)
##
## PCA needed 69 components to capture 80 percent of the variance

```

#We have a threshold limit of 0.80 for PCA to ensure 80% of variance is captured. PCA captured 80% in 69 components.

#We are adding default column from previous model to the PCA model:

```
pca_model_1$default <- bank_lasso$default
```

#Creating a train and validation sets from the values returned in PCA model:

```
set.seed(123)
```

```
pca_index <- createDataPartition(pca_model_1$default, p = 0.80, list = FALSE)
```

```
pca_train <- pca_model_1[pca_index, ]  
pca_validate <- pca_model_1[-pca_index, ]
```

#Coverting the "default" column into factor in both train and validation sets:

```
pca_train$default <- as.factor(pca_train$default)  
pca_validate$default <- as.factor(pca_validate$default)
```

#Now run the values returned from PCA in random forest model:

```
set.seed(123)
```

```
model_rf_pca <- randomForest(default ~ ., data = pca_train, mtry = 5)
```

```
print(model_rf_pca)
```

```
##
```

```
## Call:
```

```
## randomForest(formula = default ~ ., data = pca_train, mtry = 5)
```

```
##               Type of random forest: classification
```

```
##               Number of trees: 500
```

```
## No. of variables tried at each split: 5
```

```
##
```

```
##               OOB estimate of  error rate: 9.22%
```

```
## Confusion matrix:
```

```
##           0  1  class.error
```

```
## 0 58089   8 0.0001377007
```

```
## 1  5890 14 0.9976287263
```

```
pca_final <- data.frame(actual = pca_validate$default, predict(model_rf_pca,  
newdata = pca_validate, type = "prob"))
```

```
pca_final$predict <- ifelse(pca_final$X0 > 0.60, 0, 1)
```

```
pca <- confusionMatrix(as.factor(pca_final$predict),  
as.factor(pca_final$actual), positive='1')
```

```
pca
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 14510 1462
##           1   14   13
##
##           Accuracy : 0.9077
##           95% CI : (0.9032, 0.9122)
##           No Information Rate : 0.9078
##           P-Value [Acc > NIR] : 0.5178
##
##           Kappa : 0.014
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.0088136
##           Specificity : 0.9990361
##           Pos Pred Value : 0.4814815
##           Neg Pred Value : 0.9084648
##           Prevalence : 0.0921933
##           Detection Rate : 0.0008126
##           Detection Prevalence : 0.0016876
##           Balanced Accuracy : 0.5039248
##
##           'Positive' Class : 1
##
```

#Loading Test data set for predicting the defaulting customers:

```
pca_test <- read.csv("test_no_lossv3.csv")
```

#We are imputing missing values same as we did for train data set using medianimpute method:

```
test_pca_1 <- preProcess(pca_test, method = c("medianImpute"))
```

```
test_pca_process<- predict(test_pca_1, pca_test)
```

#Selecting attributes from test data set "test\_pca\_process" using coefficients returned from lasso model i.e., "new\_bank\_coef":

```
test_pca_lasso<-select(test_pca_process,
new_bank_coef[new_bank_coef!="default"])
```

#We are processing test model also in PCA to match our train model:

```
set.seed(123)
test_pca_model <- preProcess(test_pca_lasso, method = c("center", "scale",
"pca"), thresh = 0.80)
```

```
test_pca_model_1<- predict(pca_model, test_pca_lasso)
```

#Predicting the test\_pca\_model\_1 using the random forest model "model\_rf\_pca":

```
set.seed(123)
```

```
predictions_pca <-data.frame(id=pca_test$id,predict(model_rf_pca,  
test_pca_model_1, type = "prob"))
```

```
threshold <- 0.60
```

```
predictions_pca$predicted_default <- ifelse(predictions_pca$X0 > threshold,  
0, 1)
```

#Filtering the number defaulting customers that was predicting by our random forest model:

```
predictions_pca_filtered<-predictions_pca %>% filter(predicted_default == 1)  
predictions_pca_filtered
```

##		id	X0	X1	predicted_default
##	1086	34664	0.580	0.420	1
##	1342	20578	0.382	0.618	1
##	3621	90934	0.564	0.436	1
##	5105	98853	0.554	0.446	1
##	6837	4702	0.568	0.432	1
##	7645	1180	0.586	0.414	1
##	7715	13944	0.598	0.402	1
##	8647	57274	0.412	0.588	1
##	8882	99336	0.388	0.612	1
##	9128	74118	0.530	0.470	1
##	9992	21206	0.594	0.406	1
##	10263	79925	0.574	0.426	1
##	11564	40386	0.588	0.412	1
##	11616	23780	0.538	0.462	1
##	12080	2556	0.598	0.402	1
##	12892	79483	0.582	0.418	1
##	13017	71222	0.590	0.410	1
##	13140	82291	0.324	0.676	1
##	13289	64413	0.400	0.600	1
##	13352	85505	0.440	0.560	1
##	14831	86315	0.312	0.688	1
##	16841	89098	0.414	0.586	1
##	17168	2370	0.592	0.408	1
##	17987	71006	0.288	0.712	1
##	18197	89775	0.598	0.402	1
##	20030	50050	0.532	0.468	1
##	20307	85692	0.572	0.428	1
##	20333	62466	0.502	0.498	1
##	21834	1248	0.452	0.548	1
##	22427	61399	0.566	0.434	1
##	23437	56201	0.446	0.554	1

```
## 23471    835 0.474 0.526          1
## 24976 98018 0.520 0.480          1
```

#Based on the results our random forest model predicted 33 customer will default in the test data set.

#We are binding our results from our classification model to our original test dataset:

```
test_2<-pca_test
test_2$predictions <- predictions_pca$predicted_default
test_3<- test_2 %>% filter(predictions==1)
```

#2.REGRESSION MODEL: #We have classified our defaulting customers using classification above. #Now, we create a regression model, to predict loss by the defaulting customers from classification.

#Loading the train data set:

```
new_train <- read.csv("train_v3.csv")
```

#filtering all the non-defaulting customers from the train data set:

```
new_train_1 <- new_train %>% filter(loss!=0)
new_train_1$loss<- (new_train_1$loss / 100)
```

##Removing the zero-variances variables and Preprocessing the dataset by removing highly correlated and imputing missing values using “corr” and “medianimpute”:

```
zero_var_indices_1 <- nearZeroVar(new_train_1[ , -c(763)])
train_model <- new_train_1[, -zero_var_indices_1]

new_train_3 <- preProcess(train_model[ , -c(748)], method = c("medianImpute",
"corr"))


new_train_4 <- predict(new_train_3, train_model)
```

#Lasso model: we are using lasso model for variable selection for the dataset “new\_train\_4” consisting of 252 attributes:

```
set.seed(123)
x_1 <- as.matrix(new_train_4[ , -c(252)])
y_2 <- as.vector(new_train_4$loss)

model_lasso <- cv.glmnet(x_1, y_2, alpha = 1, family = "gaussian", nfolds =
10, type.measure = "mse")

plot(model_lasso)
```



```
model_lasso$lambda.min
```

```
## [1] 0.0006662975
```

#We convert coefficients returned in lasso model into dataframe:

```
# Return the coefficients for the lasso regression at the minimum Lambda value:
```

```
coef_test <- coef(model_lasso, s= "lambda.min")
```

```
#Convert the coefficient values into a data frame:
```

```
coef_test<- data.frame(name = coef_test@Dimnames[[1]][coef_test@i + 1],  
coefficient = coef_test@x)
```

```
#Removing negatives values using "abs" function:
```

```
coef_test$coefficient <- abs(coef_test$coefficient)
```

```
#Re-arranging the data frame in decreasing order:
```

```
coef_test[order(coef_test$coefficient, decreasing = TRUE), ]
```

```
##           name coefficient  
## 1  (Intercept) 1.723274e-01  
## 21          f129 1.376713e-01  
## 121         f774 1.273948e-01
```

## 119	f766 1.269790e-01
## 120	f768 1.177613e-01
## 38	f268 1.054545e-01
## 33	f229 9.832930e-02
## 78	f556 9.815358e-02
## 77	f546 5.953691e-02
## 36	f243 5.437601e-02
## 34	f238 4.915198e-02
## 15	f70 4.494179e-02
## 87	f615 4.039415e-02
## 43	f297 3.829277e-02
## 22	f130 2.363449e-02
## 29	f198 2.120761e-02
## 12	f57 2.105920e-02
## 42	f291 1.712850e-02
## 55	f402 1.638670e-02
## 85	f604 1.632679e-02
## 118	f765 1.628697e-02
## 95	f637 1.553414e-02
## 45	f316 1.462647e-02
## 92	f631 1.298546e-02
## 83	f598 1.292320e-02
## 8	f32 1.212681e-02
## 82	f589 1.154419e-02
## 91	f629 1.119956e-02
## 47	f324 1.100458e-02
## 14	f67 1.084942e-02
## 16	f71 1.020813e-02
## 13	f65 9.622098e-03
## 20	f109 7.807988e-03
## 10	f47 7.563056e-03
## 5	f13 7.314697e-03
## 37	f248 7.099032e-03
## 89	f621 6.582336e-03
## 79	f566 6.543583e-03
## 69	f509 6.434001e-03
## 48	f340 6.288034e-03
## 58	f413 6.252129e-03
## 72	f522 5.160324e-03
## 23	f140 4.389607e-03
## 30	f199 3.834597e-03
## 81	f588 3.719625e-03
## 46	f322 3.219921e-03
## 19	f93 3.125932e-03
## 18	f81 3.044201e-03
## 39	f281 3.002582e-03
## 41	f289 2.930093e-03
## 40	f288 2.827576e-03
## 28	f188 2.454033e-03
## 57	f412 2.313789e-03

## 35	f242 2.273488e-03
## 114	f739 1.966726e-03
## 64	f451 1.784469e-03
## 68	f499 1.488715e-03
## 88	f619 1.455426e-03
## 59	f422 1.419574e-03
## 66	f479 1.357014e-03
## 3	f3 1.337901e-03
## 116	f746 1.155371e-03
## 27	f150 9.941292e-04
## 111	f725 9.874954e-04
## 17	f76 9.015019e-04
## 31	f212 8.689184e-04
## 60	f436 8.067489e-04
## 24	f142 7.858284e-04
## 25	f144 7.356802e-04
## 67	f489 7.353847e-04
## 53	f384 7.256467e-04
## 113	f734 7.172839e-04
## 99	f647 7.091819e-04
## 70	f514 6.992333e-04
## 61	f442 6.894109e-04
## 9	f44 6.115054e-04
## 71	f518 4.844990e-04
## 11	f54 4.802613e-04
## 76	f536 4.538973e-04
## 86	f614 3.779220e-04
## 2	f1 3.677045e-04
## 62	f444 3.580527e-04
## 100	f648 3.432132e-04
## 44	f312 3.093357e-04
## 7	f31 2.931348e-04
## 84	f601 2.633167e-04
## 105	f664 2.230797e-04
## 51	f366 2.177829e-04
## 106	f669 2.054076e-04
## 104	f654 1.665254e-04
## 97	f639 1.635168e-04
## 108	f677 1.504512e-04
## 96	f638 1.415521e-04
## 93	f634 1.181377e-04
## 107	f674 1.180069e-04
## 98	f640 1.057873e-04
## 101	f649 8.784743e-05
## 49	f358 8.494809e-05
## 56	f403 7.949754e-05
## 52	f378 6.988383e-05
## 115	f740 6.702796e-05
## 4	f5 4.301011e-05
## 102	f651 4.213577e-05



```
## 63      f450 4.153655e-05
## 50      f361 2.601676e-05
## 117     f755 1.873274e-05
## 103     f652 1.370202e-05
## 6       f29 2.931829e-06
## 94      f636 1.392010e-06
## 80      f587 3.085592e-07
## 109     f682 2.565568e-08
## 26      f146 5.071273e-09
## 73      f523 3.885823e-09
## 74      f526 8.275708e-11
## 110     f715 1.333092e-11
## 90      f628 5.595958e-12
## 54      f401 2.974270e-12
## 32      f217 2.430969e-12
## 112     f726 1.721461e-12
## 75      f530 5.685597e-15
## 65      f472 4.191492e-38
```

*#Removing intercept columns returned from Lasso model:*

```
coef_test<- coef_test[-1, ]
```

*#Converting the data frame to a vector:*

```
coef_test<- as.vector(coef_test$name)
```


*#Adding "loss" column to the data frame:*

```
coef_test<- c(coef_test,"loss")
```

*#Selecting attributes from original data set "new\_train\_4" using coefficients returned from lasso model i.e., "coef\_test"*

```
final_model<-select(new_train_4, coef_test)
```

```
## Warning: Using an external vector in selections was deprecated in
tidyselect 1.1.0.
```

```
##  Please use `all_of()` or `any_of()` instead.
```

```
## # Was:
```

```
## data %>% select(coef_test)
```

```
##
```

```
## # Now:
```

```
## data %>% select(all_of(coef_test))
```

```
##
```

```
## See <https://tidyselect.r-lib.org/reference/faq-external-vector.html>.
```

```
## This warning is displayed once every 8 hours.
```

```
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```

##Creating a train and validation sets from lasso model dataset "final\_model":

```

set.seed(123)

bank_index_1 <- createDataPartition(final_model$loss, p = 0.80, list = FALSE)

bank_train_1 <- final_model[bank_index_1, ]
bank_validate_1 <- final_model[-bank_index_1, ]


#Creating a ridge model on the train data set from above:

x_3 <- as.matrix(bank_train_1[ , -c(121)])
y_3 <- as.vector(bank_train_1$loss)

ridge_model <- cv.glmnet(x_3, y_3, alpha = 0, family = "gaussian", nfolds =
10, type.measure = "mae")

plot(ridge_model)

```



```

ridge_model$lambda.min

## [1] 0.03637917

coef_final <- coef(ridge_model, s = "lambda.min")

```

### validating the Ridge model using “bank\_validate\_1” using “MAE” metrics:

```
x_4 <- as.matrix(bank_validate_1[, -c(121)])  
y_4 <- as.vector(bank_validate_1$loss)  
  
predicted_loss <- predict(ridge_model, s = ridge_model$lambda.min, newx =  
x_4)
```

#### ## Evaluating Performance.

```
mae <- mean(abs((predicted_loss - y_4)))  
mae_final <- cbind(y_4, predicted_loss)
```

```
print(mae)
```

```
## [1] 0.0575126
```

#Selecting attributes from original data set “test\_3” using coefficients returned from lasso model i.e., “coef\_test”

```
predict_9595 <- select(test_3, coef_test[coef_test != "loss"])
```

#Imputing missing values in updated dataset “predict\_9595”:

```
set.seed(123)  
final_preprocess <- preProcess(predict_9595, method = c("medianImpute"))
```

```
final_preprocess_1 <- predict(final_preprocess, predict_9595)
```

#Predicting loss using ridge model by defaulting customers:

```
default_loss <- as.data.frame(round(abs(predict(ridge_model, s =  
ridge_model$lambda.min, newx = as.matrix(final_preprocess_1))) * 100))
```

#Storing loss given default values into a csv file:

```
loss_given_default <- cbind.data.frame(predictions_pca_filtered,  
default_loss)
```

```
s <- left_join(predictions_pca, loss_given_default, by = 'id')
```

```
s$loss <- ifelse(s$predicted_default.x == 0, 0, s$s1)
```

```
final_predicted_file <- data.frame(id = s$id, loss = s$loss)
```

```
write.csv(final_predicted_file, "final_predicted_file.csv")
```